


Is the New Checkout Actually Better?

A/B testing an e-commerce checkout experience



The Treatment Looks Better

6,000 users randomly split · 14-day test window

USERS	CONTROL	TREATMENT
6,000	3,052	2,948

CONVERSION RATE

Control

9.50%



290 of 3,052 bought

Treatment

13.53%



399 of 2,948 bought

ABSOLUTE LIFT

+4.03

percentage points

RELATIVE LIFT

+42.4%

vs. control baseline

At first glance, the treatment looks better.
But an observed difference is not automatically a real effect.

Can We Trust the Difference?

Three checks on the same result

P-VALUE

< 0.001

Evidence against 'no difference'

Z-SCORE

4.90

Difference relative to random
variation

95% CI

[+2.42, +5.65]

Range of plausible true effects, in
points

Rates show what happened.

Statistics tell us how confident we should be.

Was the Experiment Powerful Enough?

A result you cannot detect is a result you cannot trust

MDE — 80% power **+2.23 pp**



Observed lift **+4.03 pp**



PERCENTAGE POINTS

OBSERVED LIFT > MDE

MDE is the smallest effect this experiment had a reasonable chance of detecting at 80% power. The observed lift sits well above it, so the test was sized to find an effect this large.

MDE tells you what the test could have found. It is not a significance threshold — the p-value and confidence interval remain the evidence.

FINAL DESISION MAKING

DECISION

SHIP IT

Roll out the new checkout to all users

Statistically significant?

Yes · $p < 0.001$

Effect large enough?

Yes · +4.03 pp, CI clears zero

Evidence strong enough to act?

Yes · observed lift > MDE
